

# 16.1\_pandas

May 2, 2023

## 1 Introduction to Python for Open Source Geocomputation



- Instructor: Dr. Wei Kang
- Class Location and Time: ENV 336, Mon & Wed 12:30 pm - 1:50 pm

Content:

- what is pandas?
- data processing
- data exploration
- read and save data

## 2 What is Pandas?

- Pandas is a Python library for conducting data analysis.
- First release was in 2010
- The Pandas name itself is derived from panel data, an econometrics term for multidimensional structured datasets, and a play on the phrase Python data analysis.
- Pandas provides high-level data structures and functions designed to make working with structured or tabular data intuitive and flexible.
- contains data structures and data manipulation tools designed to make data cleaning and analysis fast and convenient in Python.
- Works with **structured data**:
  - Tabular or spreadsheet-like data in which each column may be a different type (string, numeric, date, or otherwise). This includes most kinds of data commonly stored in relational databases or tab- or comma-delimited text files.
  - Multidimensional arrays (matrices).

- Multiple tables of data interrelated by key columns (what would be primary or foreign keys for a SQL user).
- Evenly or unevenly spaced time series.

## 2.1 Installation of Pandas

From a terminal:

```
pip install pandas
```

or

```
conda install pandas
```

`pandas` is included in conda installation, so our working environment should already have `pandas` installed.

```
[1]: import pandas as pd
```

## 2.2 Core of Pandas: DataFrame

- The pandas `DataFrame` is a data structure that contains **two-dimensional** data and its corresponding row and column labels.
- Pandas blends the array-computing ideas of NumPy with the kinds of data manipulation capabilities found in spreadsheets and relational databases (such as SQL).
- `DataFrames` are widely used in data science, machine learning, scientific computing, and many other data-intensive fields.
- `DataFrames` are similar to SQL tables or the spreadsheets in Excel.
- In many cases, `DataFrames` are faster, easier to use, and more powerful than tables or spreadsheets because they're an integral part of the Python and NumPy ecosystems.

### 2.2.1 What is a Pandas DataFrame?

- Represents a rectangular table of data
- Contains an ordered, named collection of columns, each of which can be a different value type (numeric, string, Boolean, etc.)
- Has both a row and column index
- Can be thought of as a dictionary of `Series` all sharing the same index.

### 2.2.2 Creating a Pandas DataFrame

- Creating from a **dictionary** of **equal-length** lists or NumPy arrays
  - key is used as the column name (string)
  - value (**equal-length** lists or NumPy arrays) is used as the records
  - The resulting `DataFrame` will have its index assigned automatically
  - The columns are placed according to the order of the keys in data

```
pd.DataFrame(dict)
```
- Creating from nested lists (sublists need to be **equal-length**) or a two-dimensional NumPy array
  - Column and row names can be specified

```
pd.DataFrame(array/nested lists, index= list, columns=list)
```

```
[2]: import numpy as np
```

```
[3]: data = {"state": ["Ohio", "Ohio", "Ohio", "Nevada", "Nevada", "Nevada"],
            "year": [2000, 2001, 2002, 2001, 2002, 2003],
            "pop": [1.5, 1.7, 3.6, 2.4, 2.9, 3.2]}
```

```
[4]: data
```

```
[4]: {'state': ['Ohio', 'Ohio', 'Ohio', 'Nevada', 'Nevada', 'Nevada'],
      'year': [2000, 2001, 2002, 2001, 2002, 2003],
      'pop': [1.5, 1.7, 3.6, 2.4, 2.9, 3.2]}
```

```
[5]: frame = pd.DataFrame(data)
      frame
```

```
[5]:
```

	state	year	pop
0	Ohio	2000	1.5
1	Ohio	2001	1.7
2	Ohio	2002	3.6
3	Nevada	2001	2.4
4	Nevada	2002	2.9
5	Nevada	2003	3.2

```
[6]: data = {"state": ["Ohio", "Ohio", "Ohio", "Nevada", "Nevada", "Nevada"],
            "pop": [1.5, 1.7, 3.6, 2.4, 2.9, 3.2],
            "year": [2000, 2001, 2002, 2001, 2002, 2003]}
      frame = pd.DataFrame(data)
      frame
```

```
[6]:
```

	state	pop	year
0	Ohio	1.5	2000
1	Ohio	1.7	2001
2	Ohio	3.6	2002
3	Nevada	2.4	2001
4	Nevada	2.9	2002
5	Nevada	3.2	2003

We can specify the order of the DataFrame's columns during the creation phase

```
[7]: frame = pd.DataFrame(data, columns=["year", "state", "pop"])
      frame
```

```
[7]:
```

	year	state	pop
0	2000	Ohio	1.5
1	2001	Ohio	1.7

```

2  2002    Ohio  3.6
3  2001  Nevada  2.4
4  2002  Nevada  2.9
5  2003  Nevada  3.2

```

```
[8]: frame = pd.DataFrame(data, columns=["year", "state"])
      frame
```

```
[8]:   year  state
0  2000   Ohio
1  2001   Ohio
2  2002   Ohio
3  2001  Nevada
4  2002  Nevada
5  2003  Nevada

```

If you pass a column that isn't contained in the dictionary, it will appear with missing values in the result:

```
[9]: frame = pd.DataFrame(data, columns=["year", "state", "pop", "poverty"])
      frame
```

```
[9]:   year  state  pop  poverty
0  2000   Ohio  1.5     NaN
1  2001   Ohio  1.7     NaN
2  2002   Ohio  3.6     NaN
3  2001  Nevada  2.4     NaN
4  2002  Nevada  2.9     NaN
5  2003  Nevada  3.2     NaN

```

```
[10]: frame.poverty
```

```
[10]: 0    NaN
      1    NaN
      2    NaN
      3    NaN
      4    NaN
      5    NaN
      Name: poverty, dtype: object

```

```
[11]: frame.poverty = 0.5
```

```
[12]: frame
```

```
[12]:   year  state  pop  poverty
0  2000   Ohio  1.5     0.5
1  2001   Ohio  1.7     0.5
2  2002   Ohio  3.6     0.5

```

```

3  2001  Nevada  2.4      0.5
4  2002  Nevada  2.9      0.5
5  2003  Nevada  3.2      0.5

```

```
[13]: type(frame)
```

```
[13]: pandas.core.frame.DataFrame
```

**Group exercise** Create a pandas DataFrame using the four array variables. The DataFrame will have four columns with names `population`, `ward`, `year` and `poverty`:

```

ward = np.tile([1,2,3,4,5], 5)
year = np.array([2000] * 5 + [2001] * 5 + [2002] * 5 + [2003] * 5 + [2004] * 5)
population = np.random.randint(5000, size=(25,))
poverty = np.random.random(size=(25,))

```

Raise your hand when you are done!

```
[14]: ward = np.tile([1,2,3,4,5], 5)
year = np.array([2000] * 5 + [2001] * 5 + [2002] * 5 + [2003] * 5 + [2004] * 5)
population = np.random.randint(5000, size=(25,))
poverty = np.random.random(size=(25,))

```

```
[15]: dict_data = {"ward":ward , "year":year, "population": population,
                  "poverty": poverty}
df = pd.DataFrame(dict_data)
df

```

```
[15]:
```

	ward	year	population	poverty
0	1	2000	1990	0.837070
1	2	2000	4139	0.763842
2	3	2000	4055	0.592783
3	4	2000	2886	0.896657
4	5	2000	2410	0.905181
5	1	2001	265	0.784550
6	2	2001	3450	0.906884
7	3	2001	3721	0.927134
8	4	2001	1247	0.001118
9	5	2001	1250	0.603424
10	1	2002	1101	0.543543
11	2	2002	3263	0.241741
12	3	2002	911	0.441993
13	4	2002	184	0.999512
14	5	2002	558	0.568797
15	1	2003	1495	0.282474
16	2	2003	2025	0.557792
17	3	2003	1208	0.466008
18	4	2003	392	0.463173

19	5	2003	2772	0.119412
20	1	2004	2453	0.912898
21	2	2004	4913	0.627135
22	3	2004	638	0.597489
23	4	2004	527	0.309552
24	5	2004	3057	0.026273

```
[16]: ward = np.tile([1,2,3,4,5], 5) # 5 wards repeat 5 times
```

```
[17]: ward
```

```
[17]: array([1, 2, 3, 4, 5, 1, 2, 3, 4, 5, 1, 2, 3, 4, 5, 1, 2, 3, 4, 5, 1, 2,
          3, 4, 5])
```

```
[18]: [1,2,3,4,5]*5
```

```
[18]: [1, 2, 3, 4, 5, 1, 2, 3, 4, 5, 1, 2, 3, 4, 5, 1, 2, 3, 4, 5, 1, 2, 3, 4, 5]
```

```
[19]: np.array([1,2,3,4,5]*5)
```

```
[19]: array([1, 2, 3, 4, 5, 1, 2, 3, 4, 5, 1, 2, 3, 4, 5, 1, 2, 3, 4, 5, 1, 2,
          3, 4, 5])
```

```
[20]: year = np.array([2000] * 5 + [2001] * 5 + [2002] * 5 + [2003] * 5 + [2004] * 5)
      ↪ # 2000, 2001, 2002, 2003, 2004 each for 5 times (5 wards)
      year
```

```
[20]: array([2000, 2000, 2000, 2000, 2000, 2001, 2001, 2001, 2001, 2001, 2002,
          2002, 2002, 2002, 2002, 2003, 2003, 2003, 2003, 2003, 2004, 2004,
          2004, 2004, 2004])
```

```
[21]: np.random.randint?
```

```
[22]: population = np.random.randint(5000, size=(25,))
```

```
[23]: population
```

```
[23]: array([2724, 4117, 4653, 4770, 2450, 1257, 2199, 2674, 3612, 3870, 1396,
          3380, 4194, 3948, 2879, 4336, 136, 221, 3441, 28, 432, 250,
          4097, 2418, 504])
```

```
[24]: poverty = np.random.random(size=(25,))
```

```
[25]: poverty
```

```
[25]: array([0.21585982, 0.67015997, 0.95536642, 0.96728885, 0.2449112 ,
          0.47637213, 0.59569255, 0.93357268, 0.26203938, 0.34815188,
```

```
0.15777429, 0.05664675, 0.66245006, 0.21250013, 0.02524779,
0.19513812, 0.27977169, 0.79045387, 0.24273107, 0.80645431,
0.97210768, 0.01545941, 0.27633935, 0.25991025, 0.67491649])
```

```
[26]: df_ward = pd.DataFrame({'population': population,
                             'ward': ward,
                             'poverty': poverty,
                             'year': year})
df_ward
```

```
[26]:
```

	population	ward	poverty	year
0	2724	1	0.215860	2000
1	4117	2	0.670160	2000
2	4653	3	0.955366	2000
3	4770	4	0.967289	2000
4	2450	5	0.244911	2000
5	1257	1	0.476372	2001
6	2199	2	0.595693	2001
7	2674	3	0.933573	2001
8	3612	4	0.262039	2001
9	3870	5	0.348152	2001
10	1396	1	0.157774	2002
11	3380	2	0.056647	2002
12	4194	3	0.662450	2002
13	3948	4	0.212500	2002
14	2879	5	0.025248	2002
15	4336	1	0.195138	2003
16	136	2	0.279772	2003
17	221	3	0.790454	2003
18	3441	4	0.242731	2003
19	28	5	0.806454	2003
20	432	1	0.972108	2004
21	250	2	0.015459	2004
22	4097	3	0.276339	2004
23	2418	4	0.259910	2004
24	504	5	0.674916	2004

Creating a pandas dataframe from a matrix/two-dimensional array

```
[27]: data = np.arange(16).reshape((4, 4))
data
```

```
[27]: array([[ 0,  1,  2,  3],
             [ 4,  5,  6,  7],
             [ 8,  9, 10, 11],
             [12, 13, 14, 15]])
```

```
[28]: df_state = pd.DataFrame(data,
                             index=["Ohio", "Colorado", "Utah", "New York"],
                             columns=["one", "two", "three", "four"])
```

```
[29]: df_state
```

```
[29]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

## 2.3 Exploring data with Pandas

```
[30]: df_ward
```

```
[30]:
```

	population	ward	poverty	year
0	2724	1	0.215860	2000
1	4117	2	0.670160	2000
2	4653	3	0.955366	2000
3	4770	4	0.967289	2000
4	2450	5	0.244911	2000
5	1257	1	0.476372	2001
6	2199	2	0.595693	2001
7	2674	3	0.933573	2001
8	3612	4	0.262039	2001
9	3870	5	0.348152	2001
10	1396	1	0.157774	2002
11	3380	2	0.056647	2002
12	4194	3	0.662450	2002
13	3948	4	0.212500	2002
14	2879	5	0.025248	2002
15	4336	1	0.195138	2003
16	136	2	0.279772	2003
17	221	3	0.790454	2003
18	3441	4	0.242731	2003
19	28	5	0.806454	2003
20	432	1	0.972108	2004
21	250	2	0.015459	2004
22	4097	3	0.276339	2004
23	2418	4	0.259910	2004
24	504	5	0.674916	2004

```
[31]: df_ward.head() # first 5 rows
```



```
[31]:
```

	population	ward	poverty	year
0	2724	1	0.215860	2000
1	4117	2	0.670160	2000
2	4653	3	0.955366	2000
3	4770	4	0.967289	2000
4	2450	5	0.244911	2000

```
[32]: df_ward.head(2) # first 2 rows
```

```
[32]:
```

	population	ward	poverty	year
0	2724	1	0.21586	2000
1	4117	2	0.67016	2000

```
[33]: df_ward.tail() # last 5 rows
```

```
[33]:
```

	population	ward	poverty	year
20	432	1	0.972108	2004
21	250	2	0.015459	2004
22	4097	3	0.276339	2004
23	2418	4	0.259910	2004
24	504	5	0.674916	2004

```
[34]: df_ward.tail(2) # last 2 rows
```

```
[34]:
```

	population	ward	poverty	year
23	2418	4	0.259910	2004
24	504	5	0.674916	2004

```
[35]: df_ward.columns
```

```
[35]: Index(['population', 'ward', 'poverty', 'year'], dtype='object')
```

```
[36]: df_ward.shape
```

```
[36]: (25, 4)
```

```
[37]: len(df_ward)
```

```
[37]: 25
```

```
[38]: df_ward.shape[0]
```

```
[38]: 25
```

```
[39]: df_ward.shape[1]
```

```
[39]: 4
```

## 2.4 Indexing DataFrame

- indexing columns
- indexing rows
  - works analogously to NumPy array indexing (integer indexing)
    - \* `iloc`: integer-based indexing.
  - you can use the index values instead of only integers
    - \* `loc`: label-based indexing

```
[40]: df_state = pd.DataFrame(data,
                             index=["Ohio", "Colorado", "Utah", "New York"],
                             columns=["one", "two", "three", "four"])

df_state
```

```
[40]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
[41]: df_state[["three", "one"]]
```

```
[41]:
```

	three	one
Ohio	2	0
Colorado	6	4
Utah	10	8
New York	14	12

```
[42]: df_state[["two"]]
```

```
[42]:
```

	two
Ohio	1
Colorado	5
Utah	9
New York	13

```
[43]: df_state["two"]
```

```
[43]:
```

Ohio	1
Colorado	5
Utah	9
New York	13

Name: two, dtype: int64

```
[44]: df_state.two
```

```
[44]:
```

Ohio	1
Colorado	5

```
Utah          9
New York      13
Name: two, dtype: int64
```

```
[45]: df_state[["three", "one"]]
```

```
[45]:
```

	three	one
Ohio	2	0
Colorado	6	4
Utah	10	8
New York	14	12

```
[46]: df_state
```

```
[46]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
[47]: df_state[:2]
```

```
[47]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7

The row selection syntax `df_state[:2]` is provided as a convenience. Passing a single element or a list to the `[]` operator selects columns.

```
[48]: df_state
```

```
[48]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
[49]: df_state[2]
```

```
-----
KeyError                                Traceback (most recent call last)
File ~/opt/anaconda3/lib/python3.9/site-packages/pandas/core/indexes/base.py:
  3802, in Index.get_loc(self, key, method, tolerance)
    3801 try:
-> 3802     return self._engine.get_loc(casted_key)
    3803 except KeyError as err:
```

```
File ~/opt/anaconda3/lib/python3.9/site-packages/pandas/_libs/index.pyx:138, in
↳ pandas._libs.index.IndexEngine.get_loc()
```

```
File ~/opt/anaconda3/lib/python3.9/site-packages/pandas/_libs/index.pyx:165, in
↳ pandas._libs.index.IndexEngine.get_loc()
```

```
File pandas/_libs/hashtable_class_helper.pxi:5745, in pandas._libs.hashtable.
↳ PyObjectHashTable.get_item()
```

```
File pandas/_libs/hashtable_class_helper.pxi:5753, in pandas._libs.hashtable.
↳ PyObjectHashTable.get_item()
```

KeyError: 2

The above exception was the direct cause of the following exception:

KeyError Traceback (most recent call last)

Cell In[49], line 1

```
----> 1 df_state[2]
```

```
File ~/opt/anaconda3/lib/python3.9/site-packages/pandas/core/frame.py:3807, in
↳ DataFrame.__getitem__(self, key)
```

```
    3805 if self.columns.nlevels > 1:
    3806     return self._getitem_multilevel(key)
-> 3807 indexer = self.columns.get_loc(key)
    3808 if is_integer(indexer):
    3809     indexer = [indexer]
```

```
File ~/opt/anaconda3/lib/python3.9/site-packages/pandas/core/indexes/base.py:
```

```
↳ 3804, in Index.get_loc(self, key, method, tolerance)
    3802     return self._engine.get_loc(casted_key)
    3803 except KeyError as err:
-> 3804     raise KeyError(key) from err
    3805 except TypeError:
    3806     # If we have a listlike key, _check_indexing_error will raise
    3807     # InvalidIndexError. Otherwise we fall through and re-raise
    3808     # the TypeError.
    3809     self._check_indexing_error(key)
```

KeyError: 2

```
[50]: df_state[:2]
```

```
[50]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7

```
[51]: df_state
```

```
[51]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
[52]: df_state[1:3]
```

```
[52]:
```

	one	two	three	four
Colorado	4	5	6	7
Utah	8	9	10	11

```
[53]: df_state[-2:]
```

```
[53]:
```

	one	two	three	four
Utah	8	9	10	11
New York	12	13	14	15

#### 2.4.1 “Row” selection on DataFrame with loc and iloc

- loc: label-based indexing
- iloc: integer-based indexing.

```
[54]: df_state
```

```
[54]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
[55]: df_state.loc["Colorado"]
```

```
[55]:
```

	one	two	three	four
one	4			
two		5		
three			6	
four				7

Name: Colorado, dtype: int64

```
[56]: df_state.loc["Utah"]
```

```
[56]:
```

	one	two	three	four
one	8			
two		9		
three			10	
four				11

Name: Utah, dtype: int64

```
[57]: df_state.iloc[1]
```

```
[57]: one      4
      two      5
      three    6
      four     7
      Name: Colorado, dtype: int64
```

```
[58]: df_state
```

```
[58]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
[59]: df_state.loc[["Utah", "Ohio"]]
```

```
[59]:
```

	one	two	three	four
Utah	8	9	10	11
Ohio	0	1	2	3

```
[60]: df_state.iloc[[2,0]]
```

```
[60]:
```

	one	two	three	four
Utah	8	9	10	11
Ohio	0	1	2	3

Filter data with conditions

```
[61]: df_state
```

```
[61]:
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

```
[62]: df_state < 9
```

```
[62]:
```

	one	two	three	four
Ohio	True	True	True	True
Colorado	True	True	True	True
Utah	True	False	False	False
New York	False	False	False	False

```
[63]: df_state[df_state < 9]
```

```
[63]:
```

	one	two	three	four
Ohio	0.0	1.0	2.0	3.0
Colorado	4.0	5.0	6.0	7.0
Utah	8.0	NaN	NaN	NaN
New York	NaN	NaN	NaN	NaN

```
[64]: df_state[df_state < 9] = 9
```

```
[65]: df_state
```

```
[65]:
```

	one	two	three	four
Ohio	9	9	9	9
Colorado	9	9	9	9
Utah	9	9	10	11
New York	12	13	14	15

```
[66]: df_state
```

```
[66]:
```

	one	two	three	four
Ohio	9	9	9	9
Colorado	9	9	9	9
Utah	9	9	10	11
New York	12	13	14	15

```
[67]: df_state.three== 10
```

```
[67]:
```

Ohio	False
Colorado	False
Utah	True
New York	False

Name: three, dtype: bool

```
[68]: df_state[df_state.three==10]
```

```
[68]:
```

	one	two	three	four
Utah	9	9	10	11

try on the other DataFrame

```
[69]: df_ward.head(2)
```

```
[69]:
```

	population	ward	poverty	year
0	2724	1	0.21586	2000
1	4117	2	0.67016	2000

```
[70]: df_ward['population']
```

```
[70]: 0      2724
      1      4117
      2      4653
      3      4770
      4      2450
      5      1257
      6      2199
      7      2674
      8      3612
      9      3870
     10      1396
     11      3380
     12      4194
     13      3948
     14      2879
     15      4336
     16        136
     17        221
     18      3441
     19         28
     20        432
     21        250
     22      4097
     23      2418
     24        504
      Name: population, dtype: int64
```

```
[71]: df_ward.population
```

```
[71]: 0      2724
      1      4117
      2      4653
      3      4770
      4      2450
      5      1257
      6      2199
      7      2674
      8      3612
      9      3870
     10      1396
     11      3380
     12      4194
     13      3948
     14      2879
     15      4336
```



```

16      136
17      221
18     3441
19       28
20     432
21     250
22    4097
23    2418
24     504
Name: population, dtype: int64

```

```
[72]: df_ward.head(2)
```

```

[72]:   population  ward  poverty  year
0         2724     1  0.21586  2000
1         4117     2  0.67016  2000

```

```
[73]: df_ward[0:4]
```

```

[73]:   population  ward  poverty  year
0         2724     1  0.215860  2000
1         4117     2  0.670160  2000
2         4653     3  0.955366  2000
3         4770     4  0.967289  2000

```

```
[74]: df_ward[-4:]
```

```

[74]:   population  ward  poverty  year
21         250     2  0.015459  2004
22        4097     3  0.276339  2004
23        2418     4  0.259910  2004
24         504     5  0.674916  2004

```

```
[75]: df_ward[df_ward.ward==2]
```

```

[75]:   population  ward  poverty  year
1         4117     2  0.670160  2000
6         2199     2  0.595693  2001
11        3380     2  0.056647  2002
16         136     2  0.279772  2003
21         250     2  0.015459  2004

```

```
[76]: df_ward[df_ward.population<1000]
```

```

[76]:   population  ward  poverty  year
16         136     2  0.279772  2003
17         221     3  0.790454  2003

```

19	28	5	0.806454	2003
20	432	1	0.972108	2004
21	250	2	0.015459	2004
24	504	5	0.674916	2004

```
[77]: df_ward[(df_ward.ward==2) & (df_ward.population < 1000)] # & binary operator to
      ↪perform and operation on lists of boolean values
```

```
[77]:      population  ward  poverty  year
16          136      2  0.279772  2003
21          250      2  0.015459  2004
```

```
[78]: (df_ward.ward==2) & (df_ward.population < 1000)
```

```
[78]: 0    False
1    False
2    False
3    False
4    False
5    False
6    False
7    False
8    False
9    False
10   False
11   False
12   False
13   False
14   False
15   False
16    True
17   False
18   False
19   False
20   False
21    True
22   False
23   False
24   False
dtype: bool
```

```
[79]: df_ward[(df_ward.ward==2) | (df_ward.population < 1000)] # | binary operator to
      ↪perform or operation on lists of boolean values
```

```
[79]:      population  ward  poverty  year
1          4117      2  0.670160  2000
6          2199      2  0.595693  2001
```

11	3380	2	0.056647	2002
16	136	2	0.279772	2003
17	221	3	0.790454	2003
19	28	5	0.806454	2003
20	432	1	0.972108	2004
21	250	2	0.015459	2004
24	504	5	0.674916	2004

```
[80]: df_ward[(~(df_ward.ward==2)) & (df_ward.population < 1000)] # not in ward 2 and
↳less than 1000 population
```

```
[80]:
```

	population	ward	poverty	year
17	221	3	0.790454	2003
19	28	5	0.806454	2003
20	432	1	0.972108	2004
24	504	5	0.674916	2004

```
[81]: df_ward[~((df_ward.ward==2) & (df_ward.population < 1000))] # not (in ward 2
↳and less than 1000 population)
```

```
[81]:
```

	population	ward	poverty	year
0	2724	1	0.215860	2000
1	4117	2	0.670160	2000
2	4653	3	0.955366	2000
3	4770	4	0.967289	2000
4	2450	5	0.244911	2000
5	1257	1	0.476372	2001
6	2199	2	0.595693	2001
7	2674	3	0.933573	2001
8	3612	4	0.262039	2001
9	3870	5	0.348152	2001
10	1396	1	0.157774	2002
11	3380	2	0.056647	2002
12	4194	3	0.662450	2002
13	3948	4	0.212500	2002
14	2879	5	0.025248	2002
15	4336	1	0.195138	2003
17	221	3	0.790454	2003
18	3441	4	0.242731	2003
19	28	5	0.806454	2003
20	432	1	0.972108	2004
22	4097	3	0.276339	2004
23	2418	4	0.259910	2004
24	504	5	0.674916	2004

### 2.4.2 Group exercise

```
ward = np.tile([1,2,3,4,5], 5)
year = np.array([2000] * 5 + [2001] * 5 + [2002] * 5 + [2003] * 5 + [2004] * 5)
population = np.random.randint(5000, size=(25,))
poverty = np.random.random(size=(25,))
df_ward = pandas.DataFrame({'population': population,
                            'ward': ward,
                            'poverty': poverty})
```

Selecting records from `df_ward` that are in ward 3, larger than 500 population, and poverty rate less than 40%

When you are done, raise your hand

```
[82]: df_ward[(df_ward.ward==3) & (df_ward.population > 500) & (df_ward.poverty<0.4)]
```

```
[82]:      population  ward  poverty  year
      22         4097      3  0.276339  2004
```

## 2.5 Creating New Columns in an existing DataFrame

```
[83]: df_ward.head()
```

```
[83]:      population  ward  poverty  year
0         2724      1  0.215860  2000
1         4117      2  0.670160  2000
2         4653      3  0.955366  2000
3         4770      4  0.967289  2000
4         2450      5  0.244911  2000
```

```
[84]: pop_pov = df_ward.population * df_ward.poverty # elementwise operation similar
      ↪ to numpy array
      pop_pov
```

```
[84]: 0      588.002144
      1      2759.048581
      2      4445.319943
      3      4613.967837
      4       600.032443
      5       598.799771
      6      1309.927918
      7      2496.373335
      8       946.486252
      9      1347.347781
     10       220.252915
     11       191.466005
     12      2778.315548
     13       838.950514
```

```

14      72.688373
15     846.118899
16      38.048950
17     174.690305
18     835.237618
19      22.580721
20     419.950516
21       3.864853
22    1132.162300
23     628.462986
24     340.157910
dtype: float64

```

```
[85]: df_ward
```

```

[85]:      population  ward  poverty  year
0         2724      1  0.215860  2000
1         4117      2  0.670160  2000
2         4653      3  0.955366  2000
3         4770      4  0.967289  2000
4         2450      5  0.244911  2000
5         1257      1  0.476372  2001
6         2199      2  0.595693  2001
7         2674      3  0.933573  2001
8         3612      4  0.262039  2001
9         3870      5  0.348152  2001
10        1396      1  0.157774  2002
11        3380      2  0.056647  2002
12        4194      3  0.662450  2002
13        3948      4  0.212500  2002
14        2879      5  0.025248  2002
15        4336      1  0.195138  2003
16         136      2  0.279772  2003
17         221      3  0.790454  2003
18        3441      4  0.242731  2003
19          28      5  0.806454  2003
20         432      1  0.972108  2004
21         250      2  0.015459  2004
22        4097      3  0.276339  2004
23        2418      4  0.259910  2004
24         504      5  0.674916  2004

```

```
[86]: df_ward['pop_pov'] = pop_pov.astype('int')
```

```
[87]: df_ward.head()
```

```
[87]:
```

	population	ward	poverty	year	pop_pov
0	2724	1	0.215860	2000	588
1	4117	2	0.670160	2000	2759
2	4653	3	0.955366	2000	4445
3	4770	4	0.967289	2000	4613
4	2450	5	0.244911	2000	600

## 2.6 Aggregation/Groupby

```
[88]: df_ward[df_ward.ward==1]
```

```
[88]:
```

	population	ward	poverty	year	pop_pov
0	2724	1	0.215860	2000	588
5	1257	1	0.476372	2001	598
10	1396	1	0.157774	2002	220
15	4336	1	0.195138	2003	846
20	432	1	0.972108	2004	419

```
[89]: df_ward.groupby(by='ward').sum()
```

```
[89]:
```

	population	poverty	year	pop_pov
ward				
1	10145	2.017252	10010	2671
2	10082	1.617730	10010	4300
3	15839	3.618182	10010	11025
4	18189	1.944470	10010	7860
5	9731	2.099682	10010	2381

```
[90]: df_ward.groupby(by='ward').sum()[['population', 'pop_pov']]
```

```
[90]:
```

	population	pop_pov
ward		
1	10145	2671
2	10082	4300
3	15839	11025
4	18189	7860
5	9731	2381

```
[91]: ward_df = df_ward.groupby(by='ward').sum()[['population', 'pop_pov']]
```

```
[92]: ward_df
```

```
[92]:
```

	population	pop_pov
ward		
1	10145	2671
2	10082	4300
3	15839	11025

4	18189	7860
5	9731	2381

```
[93]: ward_df['poverty'] = ward_df.pop_pov / ward_df.population
```

```
[94]: ward_df
```

```
[94]:
```

	population	pop_pov	poverty
ward			
1	10145	2671	0.263282
2	10082	4300	0.426503
3	15839	11025	0.696067
4	18189	7860	0.432129
5	9731	2381	0.244682

## 2.7 Joins/Merge

```
[95]: ward_df
```

```
[95]:
```

	population	pop_pov	poverty
ward			
1	10145	2671	0.263282
2	10082	4300	0.426503
3	15839	11025	0.696067
4	18189	7860	0.432129
5	9731	2381	0.244682

```
[96]: df_ward
```

```
[96]:
```

	population	ward	poverty	year	pop_pov
0	2724	1	0.215860	2000	588
1	4117	2	0.670160	2000	2759
2	4653	3	0.955366	2000	4445
3	4770	4	0.967289	2000	4613
4	2450	5	0.244911	2000	600
5	1257	1	0.476372	2001	598
6	2199	2	0.595693	2001	1309
7	2674	3	0.933573	2001	2496
8	3612	4	0.262039	2001	946
9	3870	5	0.348152	2001	1347
10	1396	1	0.157774	2002	220
11	3380	2	0.056647	2002	191
12	4194	3	0.662450	2002	2778
13	3948	4	0.212500	2002	838
14	2879	5	0.025248	2002	72
15	4336	1	0.195138	2003	846
16	136	2	0.279772	2003	38

17	221	3	0.790454	2003	174
18	3441	4	0.242731	2003	835
19	28	5	0.806454	2003	22
20	432	1	0.972108	2004	419
21	250	2	0.015459	2004	3
22	4097	3	0.276339	2004	1132
23	2418	4	0.259910	2004	628
24	504	5	0.674916	2004	340

```
[97]: df_all = df_ward.merge(ward_df, on='ward')
```

```
[98]: df_all
```

```
[98]:
```

	population_x	ward	poverty_x	year	pop_pov_x	population_y	pop_pov_y	\
0	2724	1	0.215860	2000	588	10145	2671	
1	1257	1	0.476372	2001	598	10145	2671	
2	1396	1	0.157774	2002	220	10145	2671	
3	4336	1	0.195138	2003	846	10145	2671	
4	432	1	0.972108	2004	419	10145	2671	
5	4117	2	0.670160	2000	2759	10082	4300	
6	2199	2	0.595693	2001	1309	10082	4300	
7	3380	2	0.056647	2002	191	10082	4300	
8	136	2	0.279772	2003	38	10082	4300	
9	250	2	0.015459	2004	3	10082	4300	
10	4653	3	0.955366	2000	4445	15839	11025	
11	2674	3	0.933573	2001	2496	15839	11025	
12	4194	3	0.662450	2002	2778	15839	11025	
13	221	3	0.790454	2003	174	15839	11025	
14	4097	3	0.276339	2004	1132	15839	11025	
15	4770	4	0.967289	2000	4613	18189	7860	
16	3612	4	0.262039	2001	946	18189	7860	
17	3948	4	0.212500	2002	838	18189	7860	
18	3441	4	0.242731	2003	835	18189	7860	
19	2418	4	0.259910	2004	628	18189	7860	
20	2450	5	0.244911	2000	600	9731	2381	
21	3870	5	0.348152	2001	1347	9731	2381	
22	2879	5	0.025248	2002	72	9731	2381	
23	28	5	0.806454	2003	22	9731	2381	
24	504	5	0.674916	2004	340	9731	2381	

```
poverty_y
```

0	0.263282
1	0.263282
2	0.263282
3	0.263282
4	0.263282
5	0.426503



```

6    0.426503
7    0.426503
8    0.426503
9    0.426503
10   0.696067
11   0.696067
12   0.696067
13   0.696067
14   0.696067
15   0.432129
16   0.432129
17   0.432129
18   0.432129
19   0.432129
20   0.244682
21   0.244682
22   0.244682
23   0.244682
24   0.244682

```

```
[99]: df_all = df_ward.merge(ward_df, on='ward', suffixes = ('_year', '_allyears'))
```

```
[100]: df_all
```

```
[100]:
```

	population_year	ward	poverty_year	year	pop_pov_year \
0	2724	1	0.215860	2000	588
1	1257	1	0.476372	2001	598
2	1396	1	0.157774	2002	220
3	4336	1	0.195138	2003	846
4	432	1	0.972108	2004	419
5	4117	2	0.670160	2000	2759
6	2199	2	0.595693	2001	1309
7	3380	2	0.056647	2002	191
8	136	2	0.279772	2003	38
9	250	2	0.015459	2004	3
10	4653	3	0.955366	2000	4445
11	2674	3	0.933573	2001	2496
12	4194	3	0.662450	2002	2778
13	221	3	0.790454	2003	174
14	4097	3	0.276339	2004	1132
15	4770	4	0.967289	2000	4613
16	3612	4	0.262039	2001	946
17	3948	4	0.212500	2002	838
18	3441	4	0.242731	2003	835
19	2418	4	0.259910	2004	628
20	2450	5	0.244911	2000	600
21	3870	5	0.348152	2001	1347

22	2879	5	0.025248	2002	72
23	28	5	0.806454	2003	22
24	504	5	0.674916	2004	340

	population_allyears	pop_pov_allyears	poverty_allyears
0	10145	2671	0.263282
1	10145	2671	0.263282
2	10145	2671	0.263282
3	10145	2671	0.263282
4	10145	2671	0.263282
5	10082	4300	0.426503
6	10082	4300	0.426503
7	10082	4300	0.426503
8	10082	4300	0.426503
9	10082	4300	0.426503
10	15839	11025	0.696067
11	15839	11025	0.696067
12	15839	11025	0.696067
13	15839	11025	0.696067
14	15839	11025	0.696067
15	18189	7860	0.432129
16	18189	7860	0.432129
17	18189	7860	0.432129
18	18189	7860	0.432129
19	18189	7860	0.432129
20	9731	2381	0.244682
21	9731	2381	0.244682
22	9731	2381	0.244682
23	9731	2381	0.244682
24	9731	2381	0.244682

```
[101]: df_all[df_all.poverty_year > df_all.poverty_allyears]
```

```
[101]:
```

	population_year	ward	poverty_year	year	pop_pov_year \
1	1257	1	0.476372	2001	598
4	432	1	0.972108	2004	419
5	4117	2	0.670160	2000	2759
6	2199	2	0.595693	2001	1309
10	4653	3	0.955366	2000	4445
11	2674	3	0.933573	2001	2496
13	221	3	0.790454	2003	174
15	4770	4	0.967289	2000	4613
20	2450	5	0.244911	2000	600
21	3870	5	0.348152	2001	1347
23	28	5	0.806454	2003	22
24	504	5	0.674916	2004	340

	population_allyears	pop_pov_allyears	poverty_allyears
1	10145	2671	0.263282
4	10145	2671	0.263282
5	10082	4300	0.426503
6	10082	4300	0.426503
10	15839	11025	0.696067
11	15839	11025	0.696067
13	15839	11025	0.696067
15	18189	7860	0.432129
20	9731	2381	0.244682
21	9731	2381	0.244682
23	9731	2381	0.244682
24	9731	2381	0.244682

Which ward has the highest average poverty rate?

```
[102]: df_all.poverty_allyears.idxmax()
```

```
[102]: 10
```

```
[103]: df_all.loc[df_all['poverty_allyears'].idxmax()]
```

```
[103]: population_year      4653.000000
ward                    3.000000
poverty_year           0.955366
year                  2000.000000
pop_pov_year          4445.000000
population_allyears    15839.000000
pop_pov_allyears       11025.000000
poverty_allyears        0.696067
Name: 10, dtype: float64
```

Which ward in which year has the lowest poverty rate?

```
[104]: df_all.poverty_year.idxmin()
```

```
[104]: 9
```

```
[105]: df_all.loc[df_all['poverty_year'].idxmin()]
```

```
[105]: population_year      250.000000
ward                    2.000000
poverty_year           0.015459
year                  2004.000000
pop_pov_year          3.000000
population_allyears    10082.000000
pop_pov_allyears       4300.000000
poverty_allyears        0.426503
```

Name: 9, dtype: float64

## 2.8 Reading and Writing Data with Pandas

- Pandas features a number of functions for reading tabular data as a `DataFrame` object.
- Works with many different data formats
- Works with different data source:
  - reading text files and other more efficient on-disk formats
  - loading data from databases
  - interacting with network sources like web APIs

### 2.8.1 An example with working with csv files

- `read_csv` function: Load delimited data from a file, URL, or file-like object; use comma as default delimiter
  - A long list of optional arguments to deal with messy data in the real world
- `to_csv` method (associated with a `DataFrame` instance): Writing to a csv file

```
[106]: df1 = pd.read_csv("ex1.csv")
df1
```

```
[106]:   a  b  c  d message
0  1  2  3  4  hello
1  5  6  7  8  world
2  9 10 11 12   foo
```

If only the path is supplied, the first row of the file will be used as the header (column names) of the `DataFrame` object and column names are inferred from the first line of the file.

```
[107]: df2 = pd.read_csv("ex1.csv", header=None)
df2
```

```
[107]:   0  1  2  3  4
0  a  b  c  d message
1  1  2  3  4  hello
2  5  6  7  8  world
3  9 10 11 12   foo
```

If `header=None`, integer index starting from 0 will be used as column names.

```
[108]: df3 = pd.read_csv("ex1.csv", names=["col1", "col2", "col3", "col4", "col5"])
df3
```

```
[108]:   col1 col2 col3 col4  col5
0    a    b    c    d message
1    1    2    3    4  hello
2    5    6    7    8  world
3    9   10   11   12   foo
```

We can pass a list of column names to the argument `names`

```
[109]: df4 = pd.read_csv("ex1.csv", index_col="message")
df4
```

```
[109]:      a  b  c  d
message
hello    1  2  3  4
world    5  6  7  8
foo      9 10 11 12
```

We can specify the column name/index in the argument `index_col` as the row labels of the DataFrame

```
[110]: df4 = pd.read_csv("ex1.csv", index_col=4)
df4
```

```
[110]:      a  b  c  d
message
hello    1  2  3  4
world    5  6  7  8
foo      9 10 11 12
```

```
[111]: df5 = pd.read_csv("ex1.csv", skiprows=[1,2])
df5
```

```
[111]:   a  b  c  d message
0  9 10 11 12      foo
```

Argument `skiprows`: Line numbers to skip (0-indexed) or number of lines to skip (int) at the start of the file.

```
[112]: df6 = pd.read_csv("ex1.csv", skiprows=2)
df6
```

```
[112]:   5  6  7  8 world
0  9 10 11 12      foo
```

## Dealing with missing values

- To control which values are parsed as missing values (which are signified by NaN), specify a string in `na_values`.
- If you specify a list of strings, then all values in it are considered to be missing values.
- If you specify a number (a float, like 5.0 or an integer like 5), the corresponding equivalent values will also imply a missing value (in this case effectively [5.0, 5] are recognized as NaN).

```
[113]: df_ex5 = pd.read_csv("ex5.csv")
df_ex5
```

```
[113]: something a    b    c    d message
0         one  1    2    3.0    4      NaN
1         two  5    6    NaN    8    world
2        three  9   10   11.0   12      foo
```

```
[114]: df_ex5 = pd.read_csv("ex5.csv", na_values=["one", 1])
df_ex5
```

```
[114]: something    a    b    c    d message
0         NaN  NaN    2    3.0    4      NaN
1         two  5.0    6    NaN    8    world
2        three  9.0   10   11.0   12      foo
```

```
[115]: df_ex5.dropna() #Drop the rows where at least one element is missing.
```

```
[115]: something    a    b    c    d message
2        three  9.0   10   11.0   12      foo
```

```
[116]: df_ex5.dropna(axis='columns') # Drop the columns where at least one element is
↳missing.
```

```
[116]:      b    d
0     2    4
1     6    8
2    10   12
```

```
[117]: df_ex5.dropna(subset=["something"]) #Define in which columns to look for
↳missing values.
```

```
[117]: something    a    b    c    d message
1         two  5.0    6    NaN    8    world
2        three  9.0   10   11.0   12      foo
```

Save a Dataframe to a csv file

```
[118]: df4.to_csv("data/output1.csv")
```

Read panda's [documentation](#) to better understand the functionality of pandas's `read_csv` function.

### 3 Further readings

- [Python for Data Analysis, 3E](#), by Wes McKinney